

Johns Hopkins Engineering

Computer Vision

Corner Detection and Matching



JOHNS HOPKINS
WHITING SCHOOL
of ENGINEERING

Corner Detection and Matching

- Topics:
 - Motivating Feature Detection
 - Harris Corner Detector Theory
 - Matching With Patches As Features

Motivation: Build a Panorama



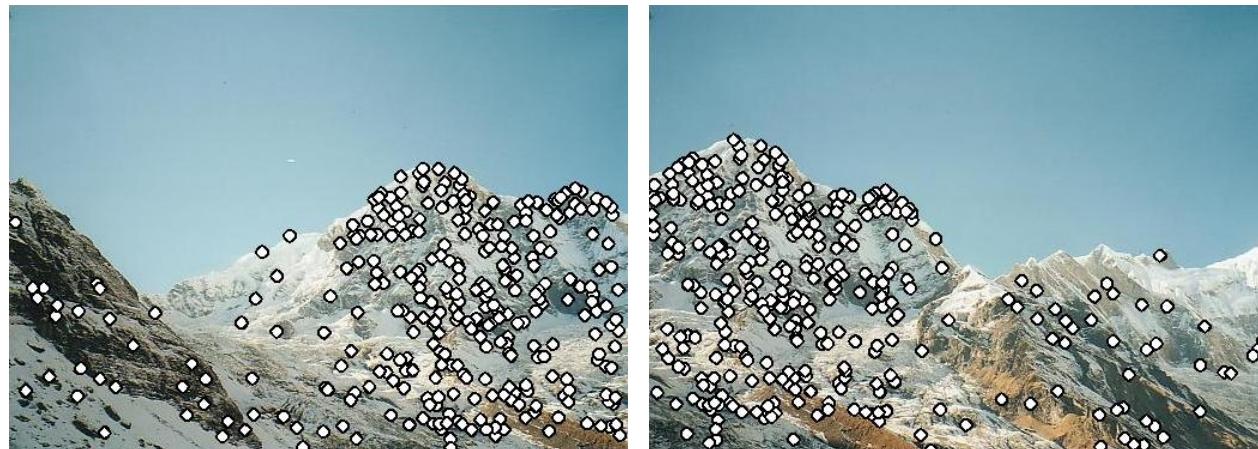
How do we build a panorama?

- We need to match (align) images



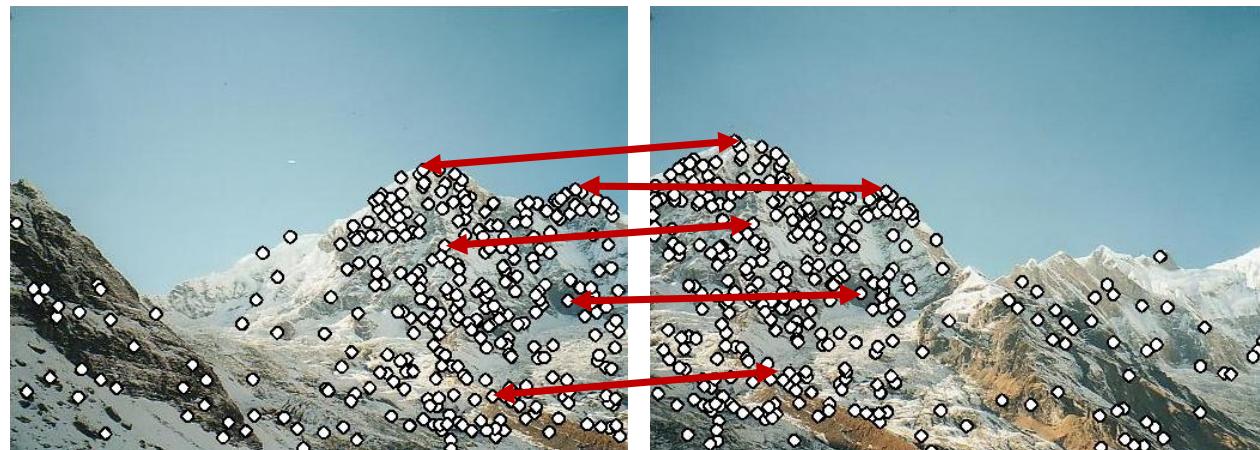
Matching with Features (1)

- Detect feature points in both images



Matching with Features (2)

- Detect feature points in both images
- Find corresponding pairs



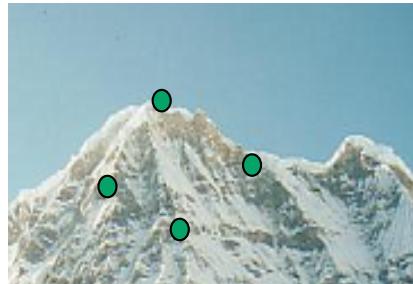
Matching with Features (3)

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images



Matching with Features (4)

- Problem 1:
 - Detect the *same* point *independently* in both images

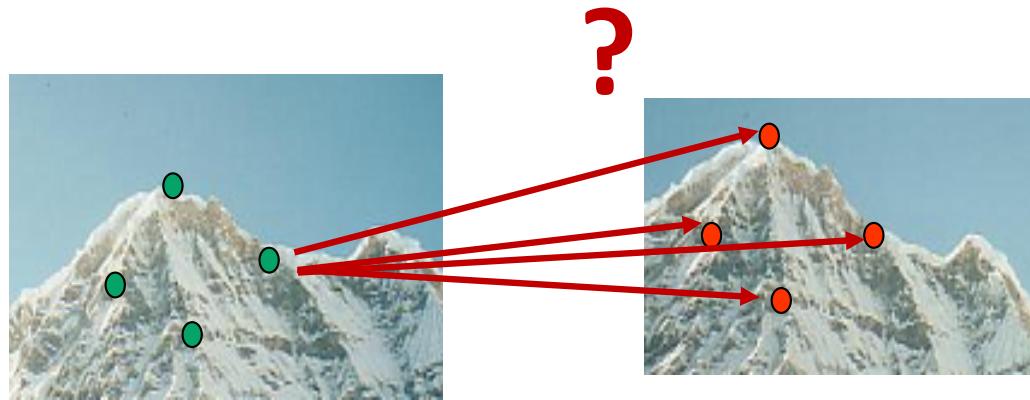


no chance to match!

We need a repeatable detector

Matching with Features (5)

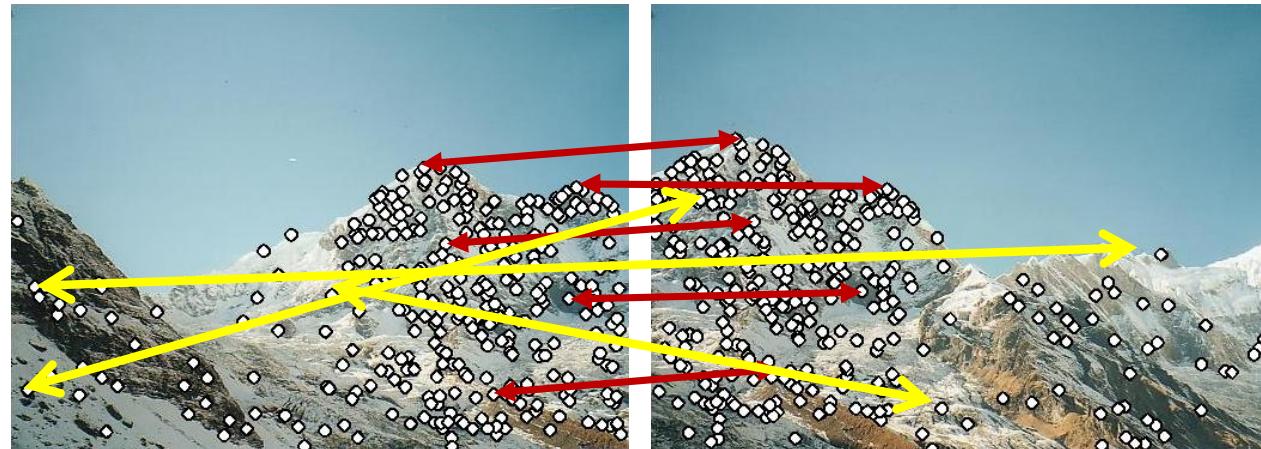
- Problem 2:
 - For each point correctly recognize the corresponding one



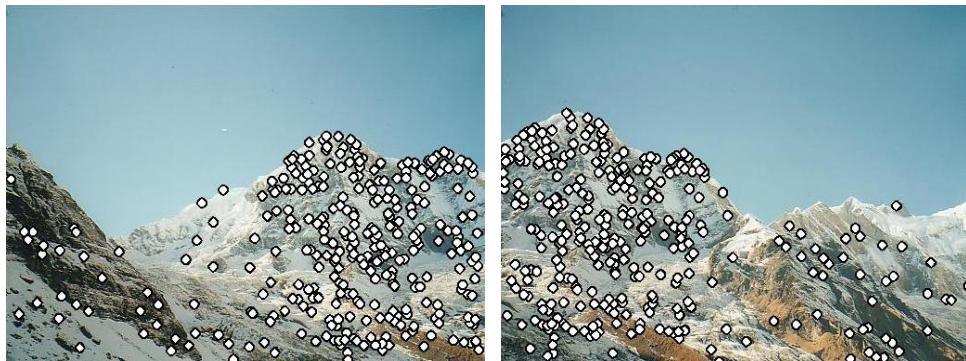
We need a reliable and distinctive descriptor

Matching with Features

- Problem 3:
 - Need to estimate transformation between images, despite erroneous correspondences.



Characteristics of Good Features

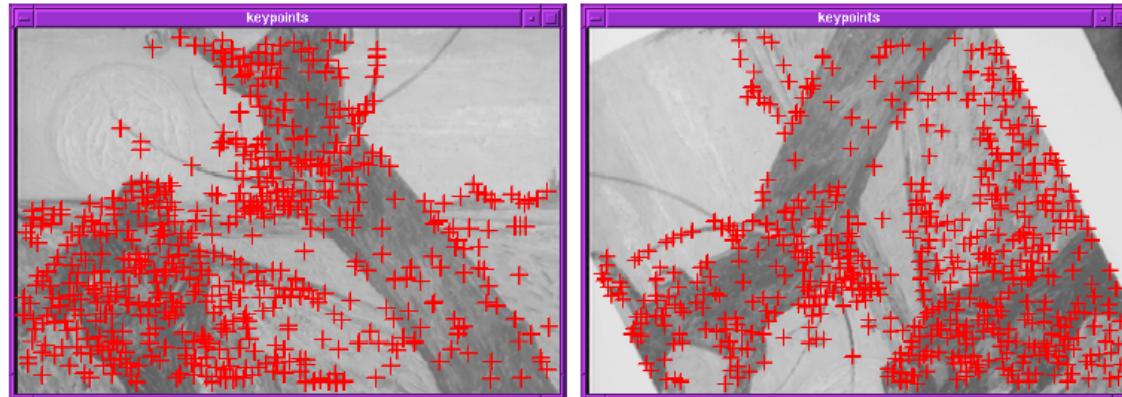


- Repeatability
 - The same feature can be found in several images despite geometric and photometric transformations
- Saliency
 - Each feature has a distinctive description
- Compactness and efficiency
 - Many fewer features than image pixels
- Locality
 - A feature occupies a relatively small area of the image; robust to clutter and occlusion

Applications

- Feature points are used for:
 - Motion tracking
 - Image alignment
 - 3D reconstruction
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation

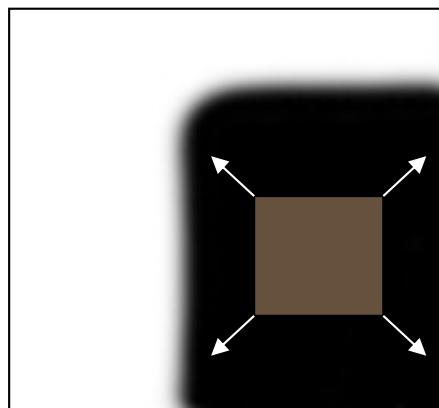
Corners



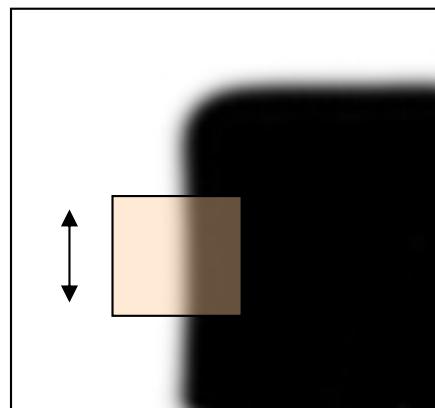
- Key property: in the region around a corner, image gradient has two or more dominant directions
- Corners are repeatable and distinctive

Corner Detection: Basic Idea

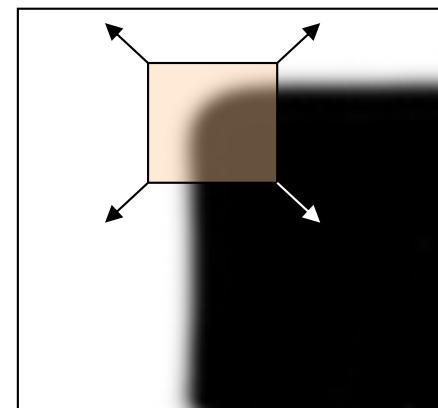
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

Image Gradients

I

Flat Region



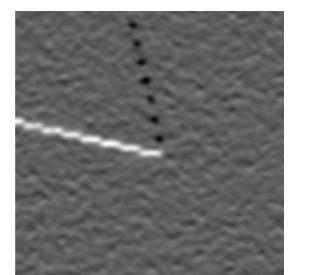
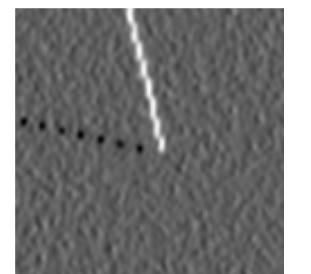
$$I_x = \frac{\partial I}{\partial x}$$

Edge Region



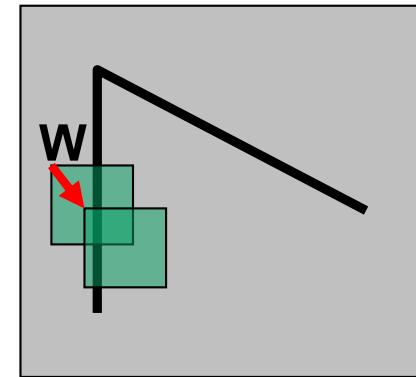
$$I_y = \frac{\partial I}{\partial y}$$

Corner Region



Feature Detection: The Math

- Consider shifting the window \mathbf{W} by (u, v)
 - How do the pixels in \mathbf{W} change?
 - Compare each pixel before and after by summing up the squared differences (SSD)
 - This defines an SSD “error” of $E(u, v)$:



$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

Small Motion Assumption

- Taylor Series expansion of I :

$$I(x + u, y + v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \dots \text{higher order terms...}$$

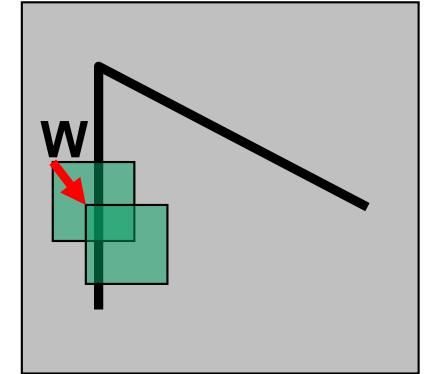
- If the motion (u, v) is small, then first order approx. is good

$$\begin{aligned} I(x + u, y + v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned} \quad \text{shorthand: } I_x = \frac{\partial I}{\partial x}$$

- Plugging this into the formula on the previous slide...

Feature Detection: The Math (2)

- Consider shifting the window \mathbf{W} by (u, v)
 - How do the pixels in \mathbf{W} change?
 - Compare each pixel before and after by summing up the squared differences (SSD)
 - This defines an SSD “error” of $E(u, v)$:

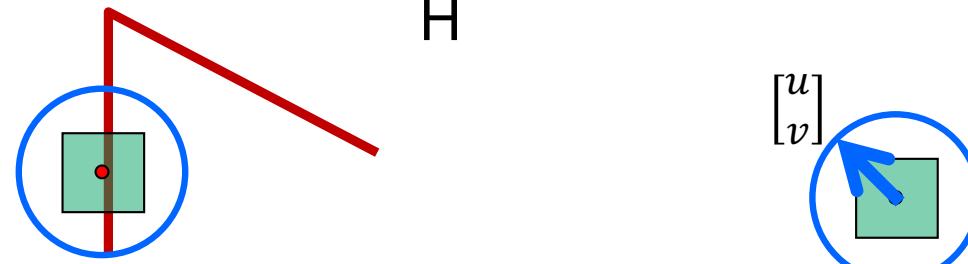


$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} \left[I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} - I(x, y) \right]^2 \\ &\approx \sum_{(x,y) \in W} \left[[I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right]^2 \end{aligned}$$

Feature Detection: The Math (3)

- This can be rewritten:

$$E(u, v) = [u \quad v] \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



For the example above

- You can move the center of the blue window to anywhere on the blue unit circle
- Which directions will result in the largest and smallest E values?
- We can find these directions by looking at the eigenvectors of H

Eigenvalue/Eigenvector Review

The **eigenvectors** of a matrix \mathbf{A} are the vectors \mathbf{x} that satisfy: $A\mathbf{x} = \lambda\mathbf{x}$

The scalar λ is the **eigenvalue** corresponding to \mathbf{x}

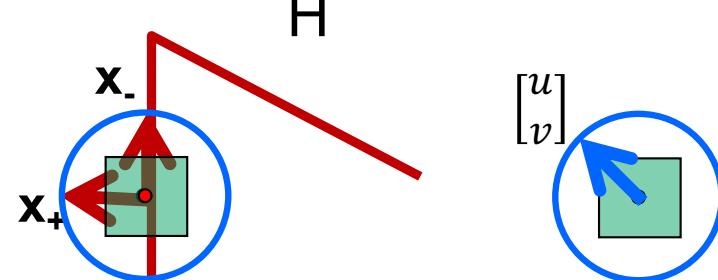
- The eigenvalues are found by solving: $\det(A - \lambda I) = 0$
- In our case, $\mathbf{A} = \mathbf{H}$ is a 2×2 matrix, so we have: $\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$
- The solution: $\lambda_{\pm} = \frac{1}{2} \left[(h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$

Once you know λ , you find \mathbf{x} by solving $\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$

Feature Detection: The Math (4)

- This can be rewritten:

$$E(u, v) = [u \quad v] \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



Eigenvalues and eigenvectors of H

- Define shifts with the smallest and largest change (E value)
- x_+ = direction of **largest** increase in E .
- λ_+ = amount of increase in direction x_+
- x_- = direction of **smallest** increase in E .
- λ_- = amount of increase in direction x_-

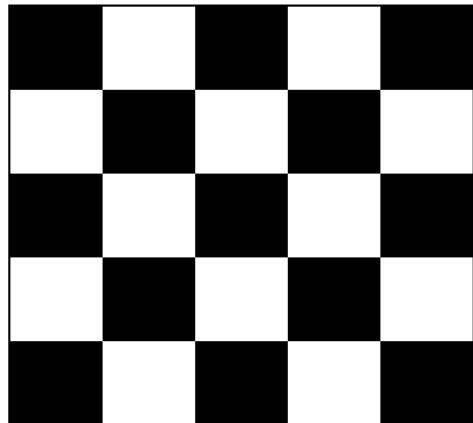
$$\begin{aligned} Hx_+ &= \lambda_+ x_+ \\ Hx_- &= \lambda_- x_- \end{aligned}$$

Feature Detection: The Math (5)

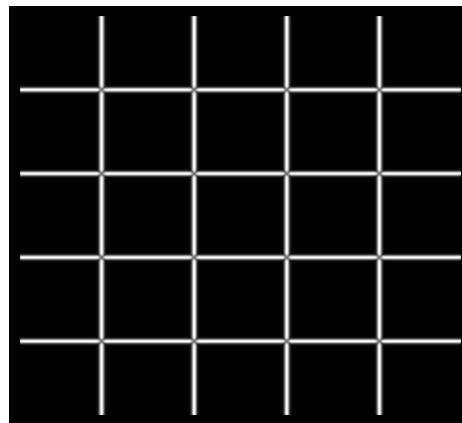
- How are λ_+ , \mathbf{x}_+ , λ_- , and \mathbf{x}_- relevant for feature detection?
 - What's our feature scoring function?

Feature Detection: The Math (6)

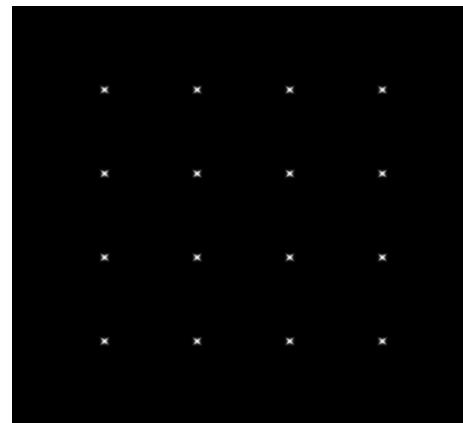
- How are λ_+ , \mathbf{x}_+ , λ_- , and \mathbf{x}_- relevant for feature detection?
 - What's our feature scoring function?
- Want $E(u,v)$ to be **large** for small shifts in **all** directions
 - The *minimum* of $E(u,v)$ should be large, over all unit vectors $[u \ v]$
 - This minimum is given by the smaller eigenvalue (λ_-) of H



|



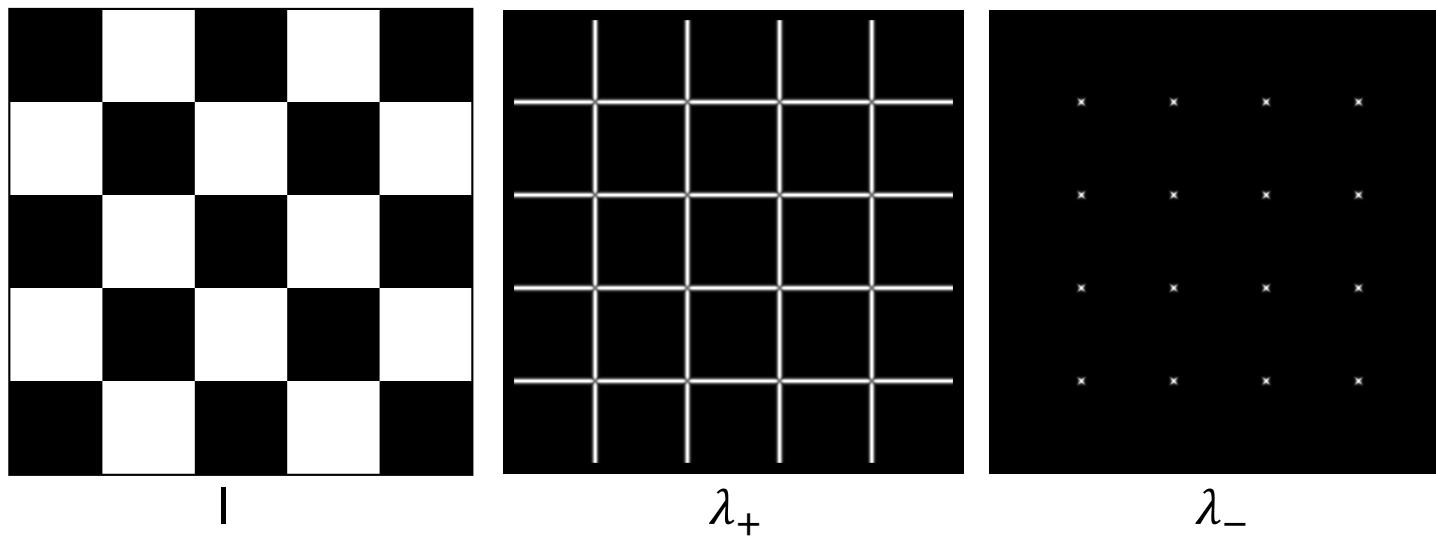
λ_+



λ_-

Feature Detection: The Math (7)

- Here's what you do
 - Compute the gradient at each point in the image
 - Create the H matrix from the entries in the gradient
 - Compute the eigenvalues.
 - Find points with large response ($\lambda_- > \text{threshold}$)
 - Choose those points where λ_- is a local maximum as features



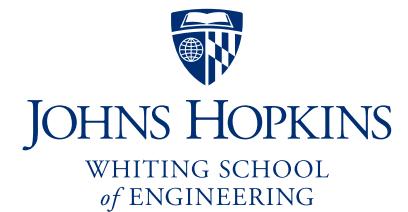
Summary

- **Corners:** Are one way to think about stable feature locations in images
- Essential concepts in this lecture:
 - What we mean by corners as a distribution of gradients
 - Formulating the corner detection problem
 - Deriving a measure for corners

Johns Hopkins Engineering

Computer Vision

Corner Detection and Matching

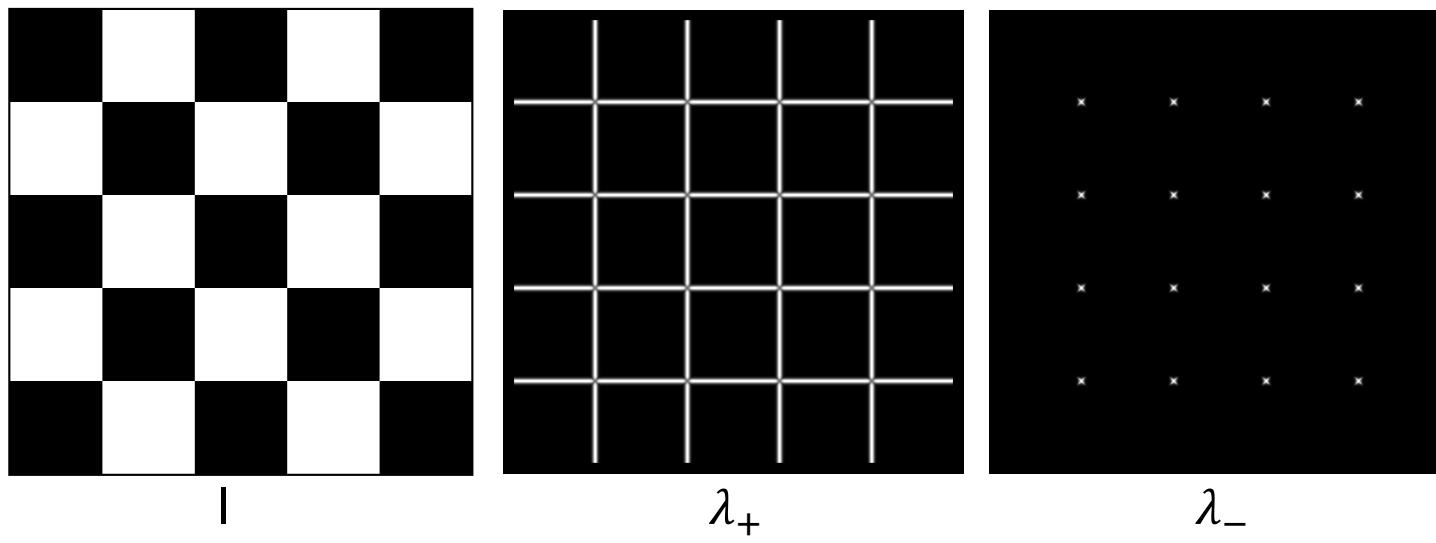


Corner Detection and Matching

- Topics:
 - Harris Corner Detector Properties
 - Harris Corner Detector Examples

Feature Detection: The Math (7)

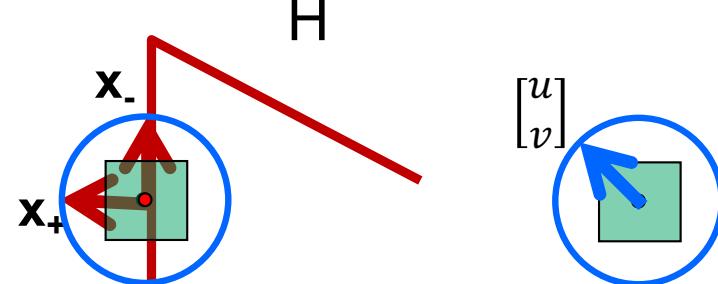
- Here's what you do
 - Compute the gradient at each point in the image
 - Create the H matrix from the entries in the gradient
 - Compute the eigenvalues.
 - Find points with large response ($\lambda_- > \text{threshold}$)
 - Choose those points where λ_- is a local maximum as features



Feature Detection: The Math (4)

- This can be rewritten:

$$E(u, v) = [u \quad v] \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

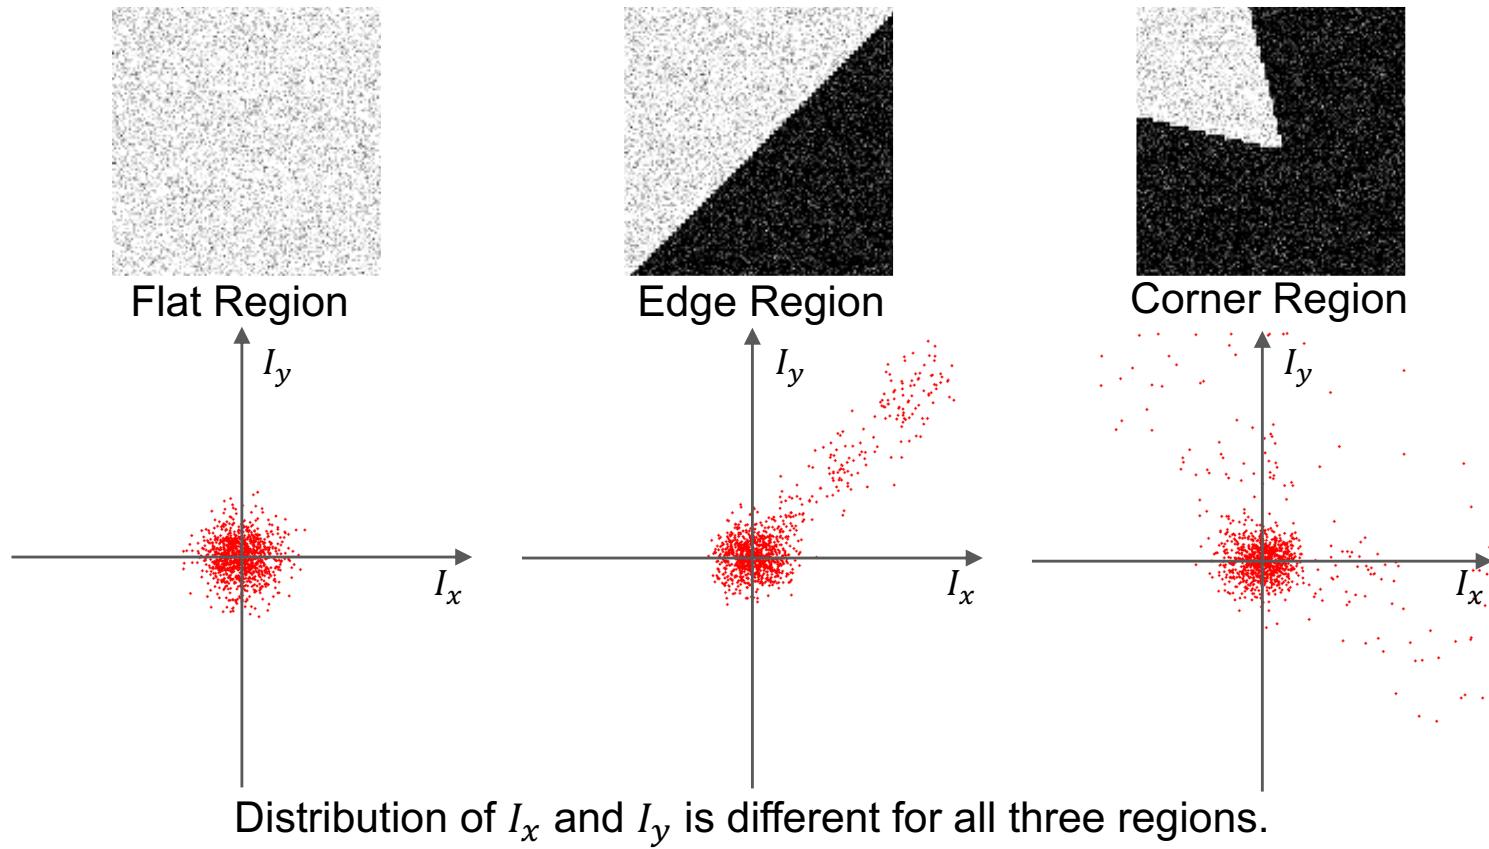


Eigenvalues and eigenvectors of H

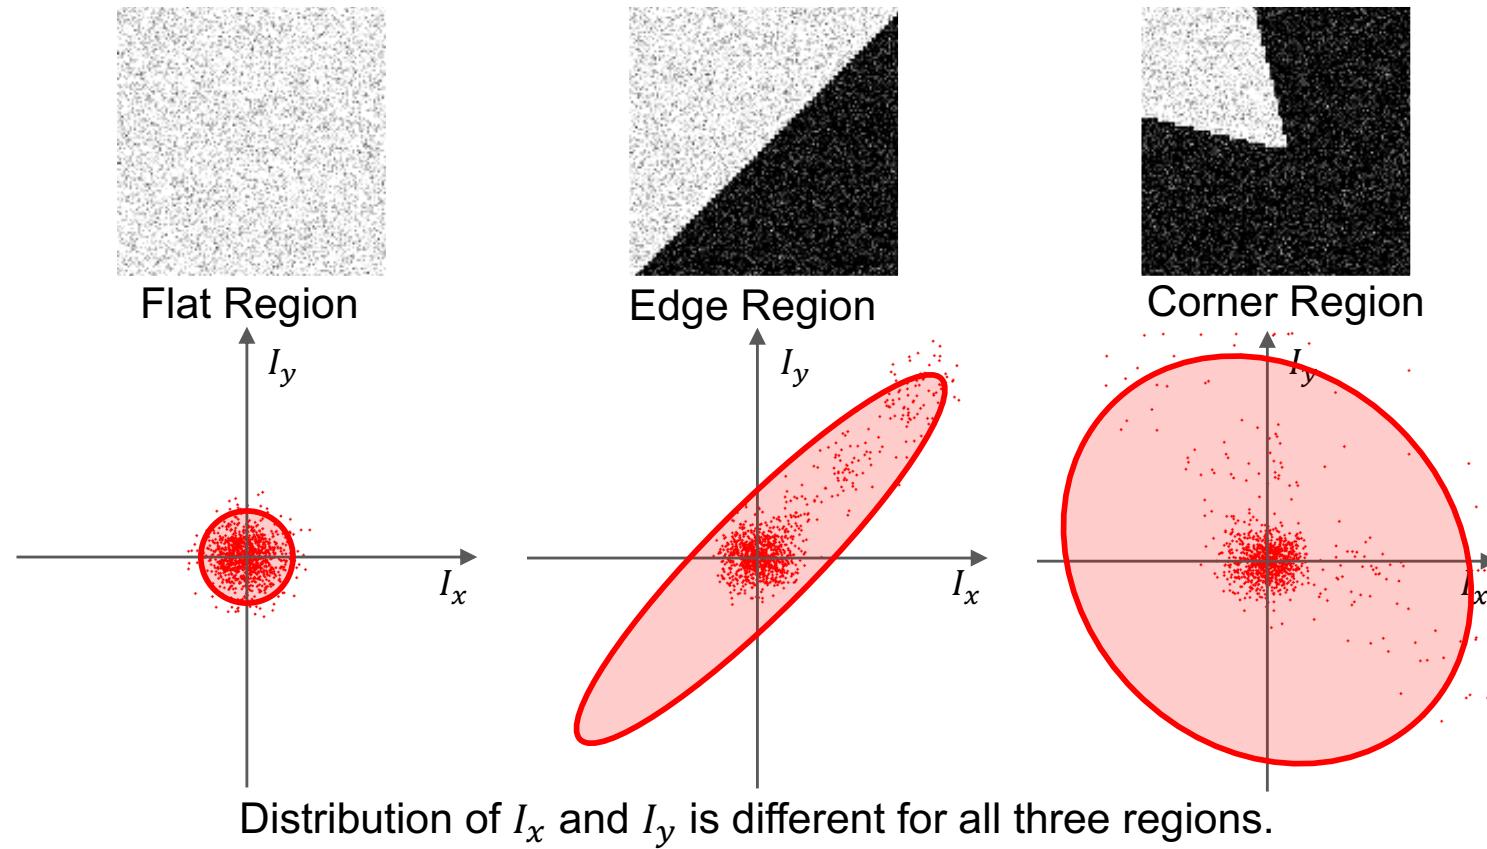
- Define shifts with the smallest and largest change (E value)
- x_+ = direction of **largest** increase in E .
- λ_+ = amount of increase in direction x_+
- x_- = direction of **smallest** increase in E .
- λ_- = amount of increase in direction x_-

$$\begin{aligned} Hx_+ &= \lambda_+ x_+ \\ Hx_- &= \lambda_- x_- \end{aligned}$$

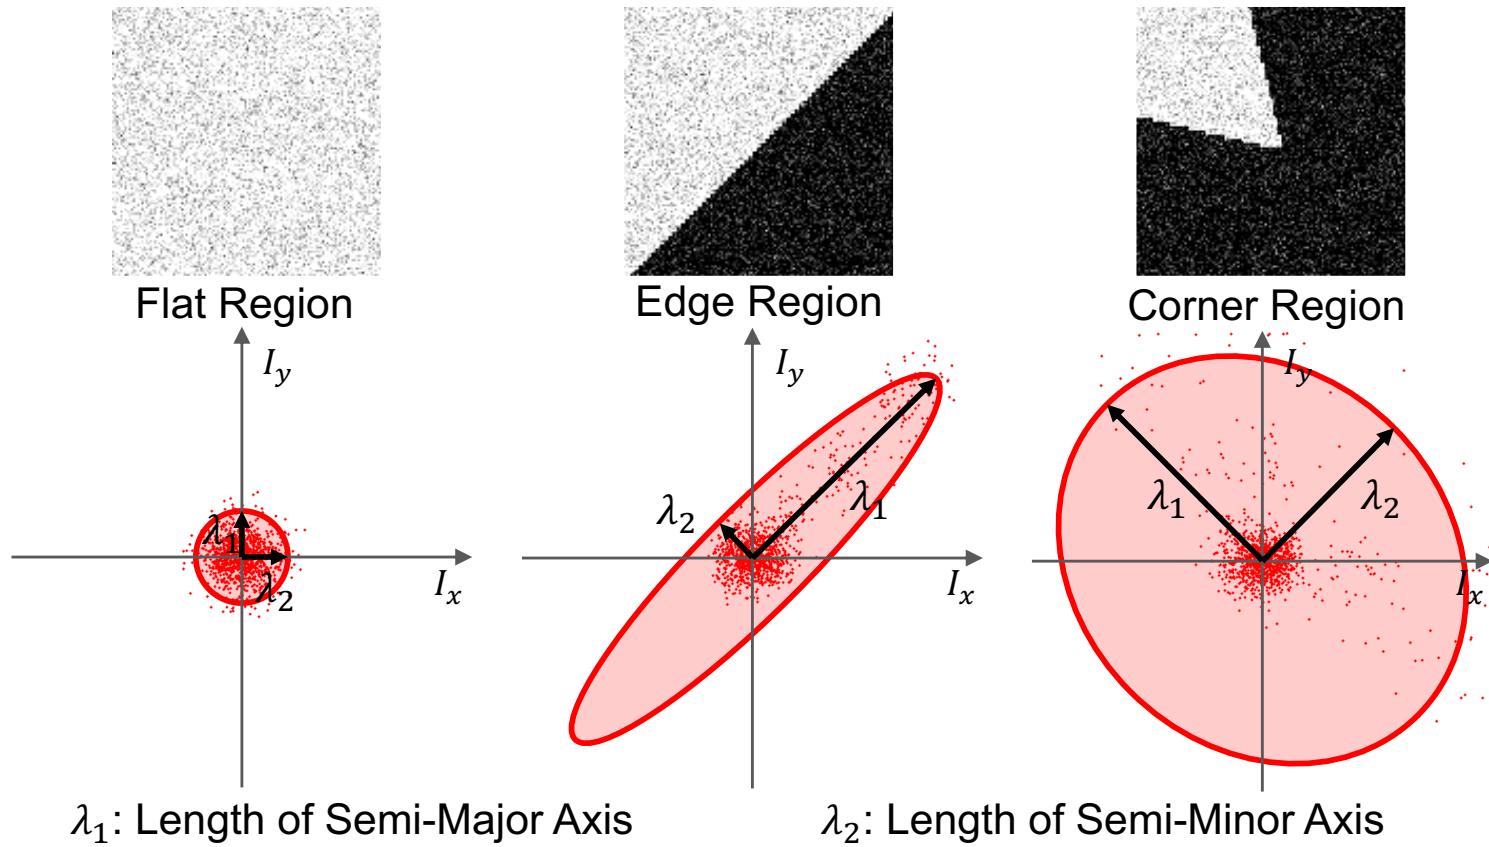
Distribution of Image Gradients



Fitting Elliptical Disk to Distribution



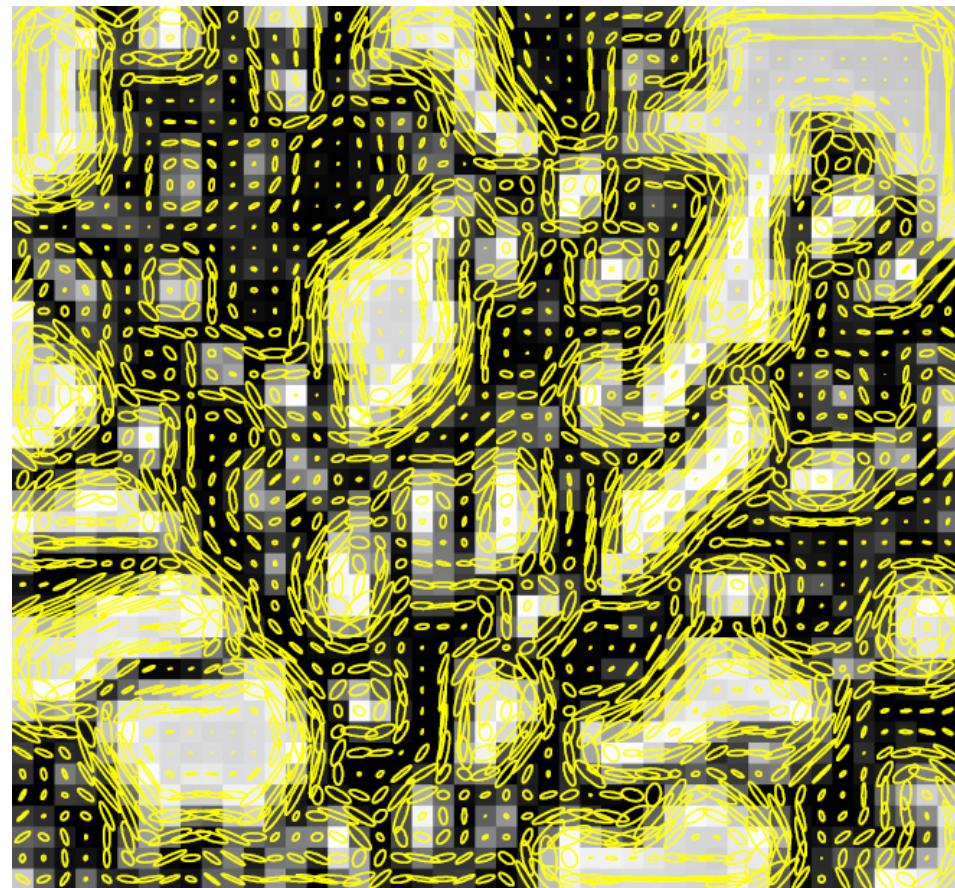
Fitting Elliptical Disk to Distribution (cont.)



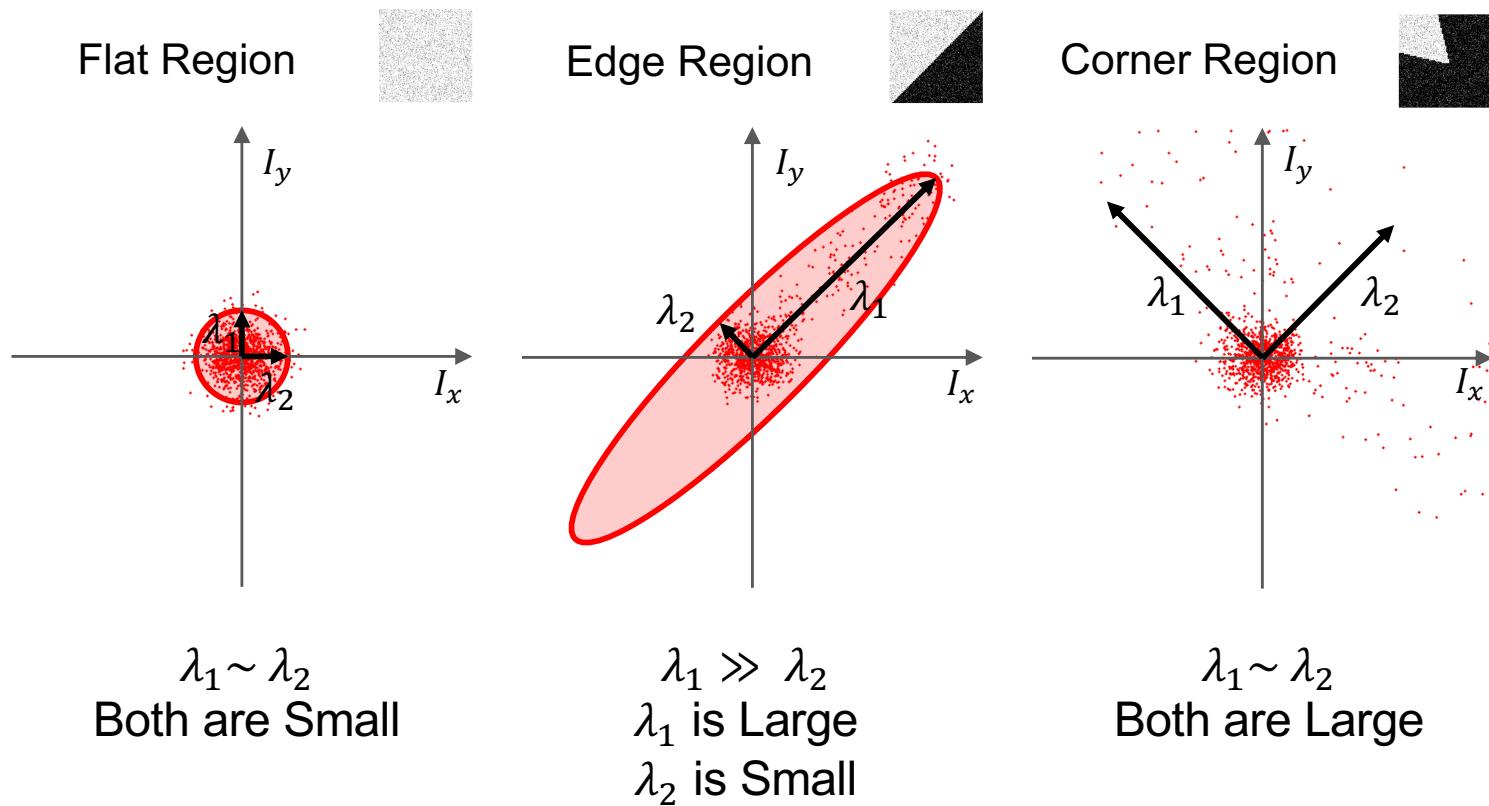
Visualization of Second Moments



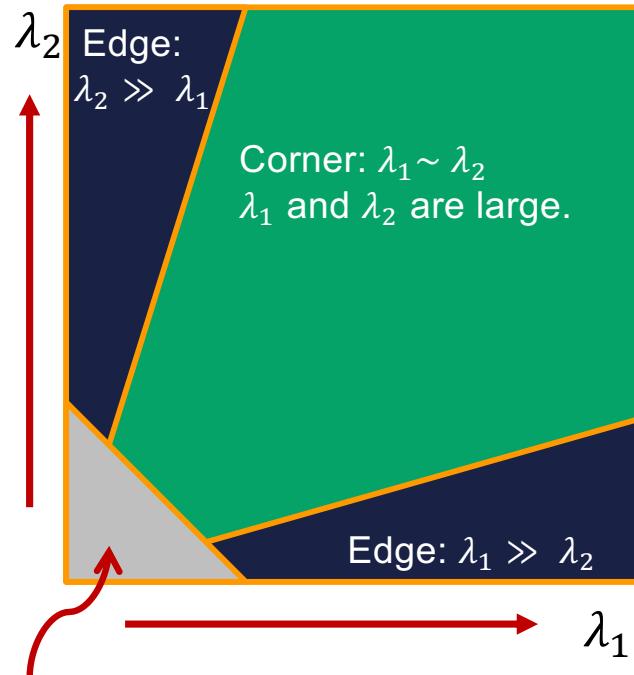
Visualization of Second Moments



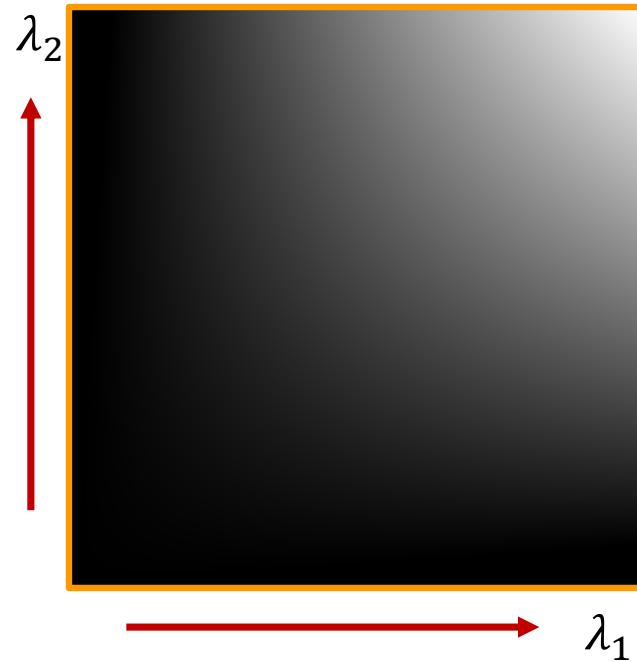
Interpretation of λ_1 and λ_2



Harris Corner Response Function



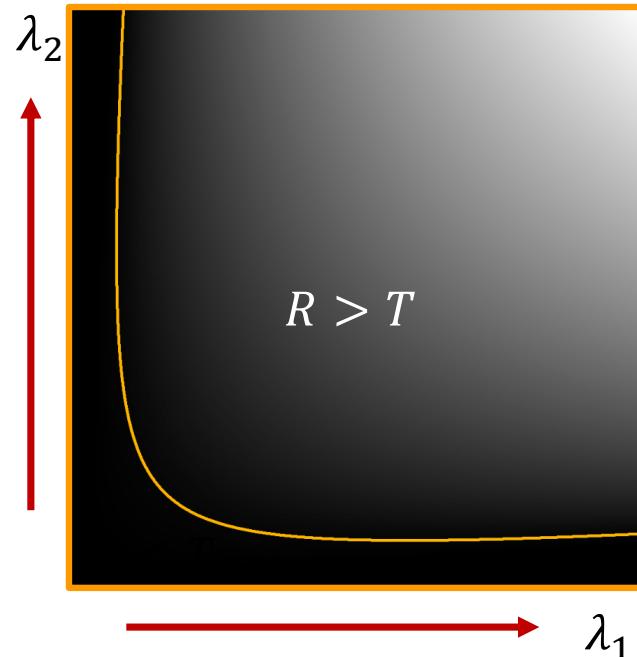
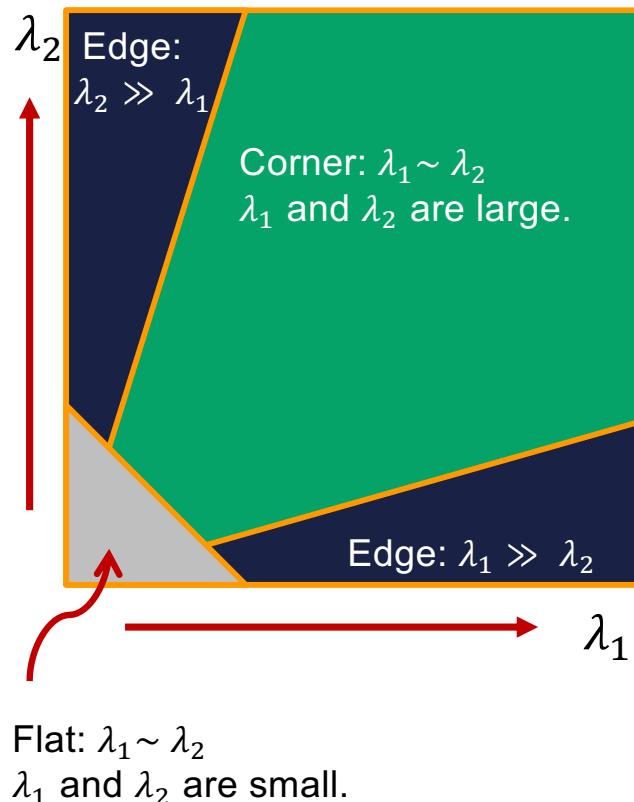
Flat: $\lambda_1 \sim \lambda_2$
 λ_1 and λ_2 are
small.



$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$

where: $0.04 \leq k \leq 0.06$
(Designed Empirically)

Harris Corner Response Function (cont.)



$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$

where: $0.04 \leq k \leq 0.06$
(Designed Empirically)

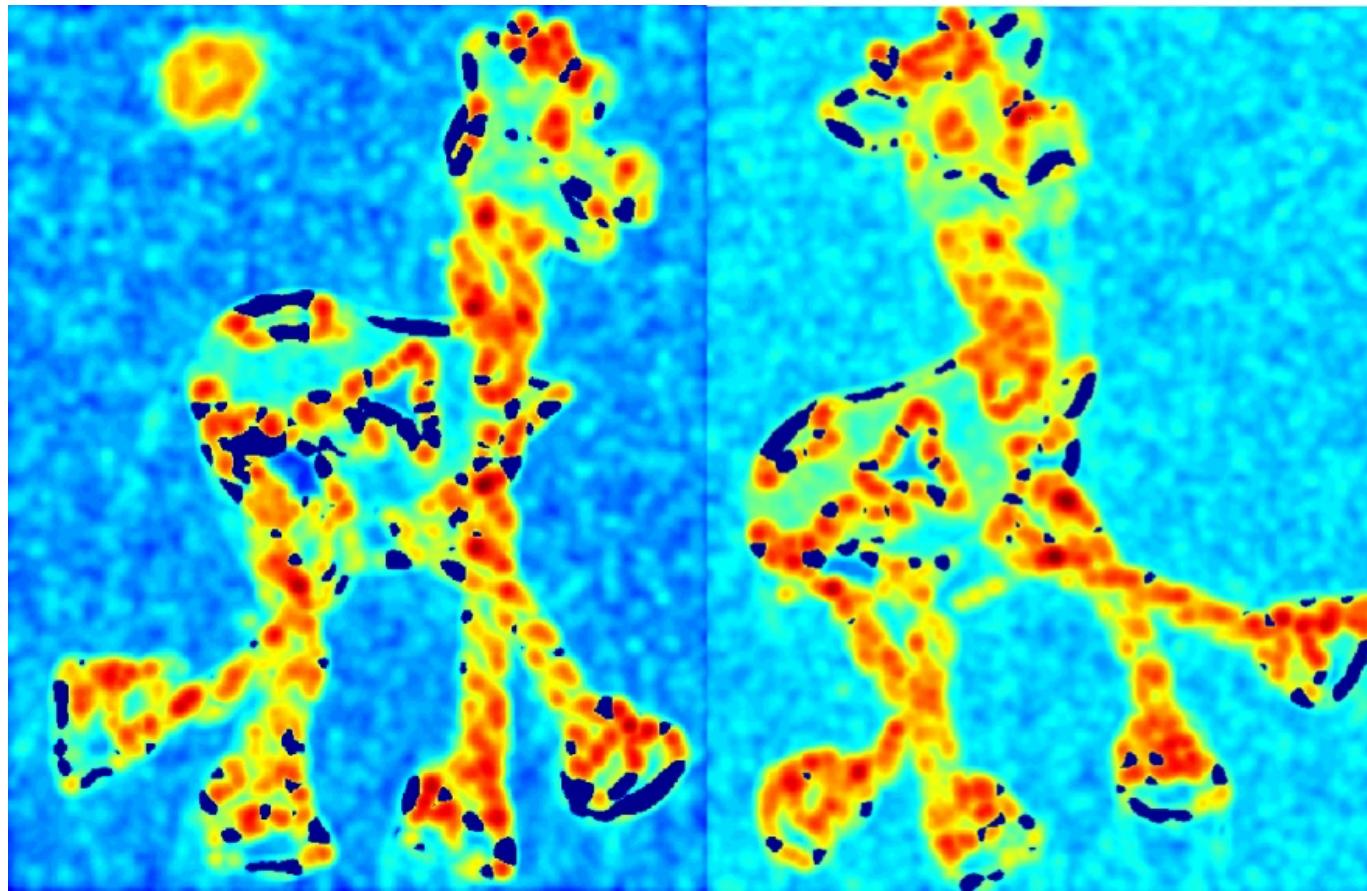
Harris Detector: Steps

1. Compute Gaussian derivatives at each pixel
2. Compute second moment matrix H in a Gaussian window around each pixel
3. Compute corner response function R
4. Threshold R
5. Find local maxima of response function (nonmaximum suppression)

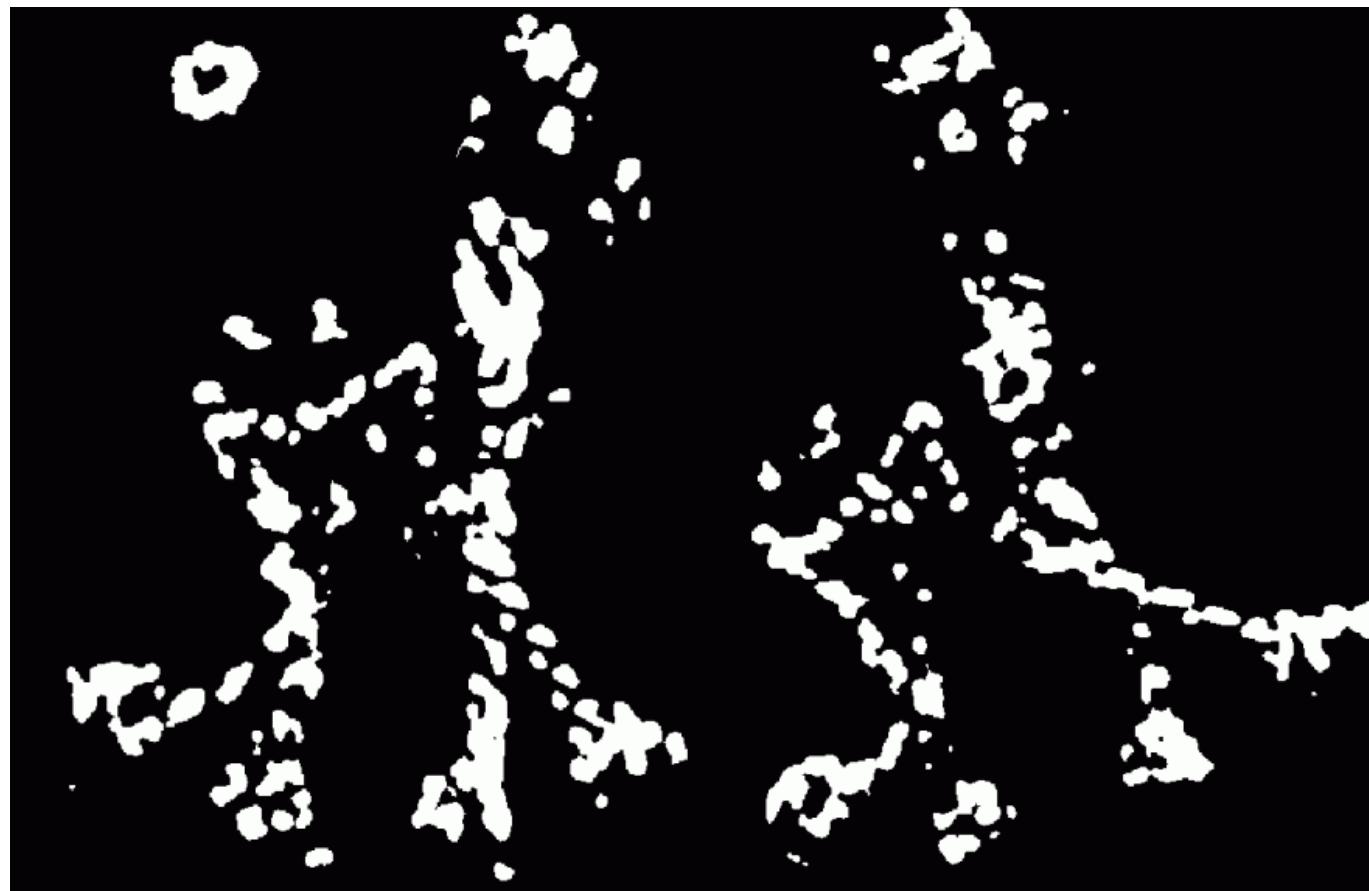
Harris Detector: Steps (cont.)



Harris: Corner Response



Harris: Threshold Response



Harris: Local Maxima



Harris Corner Locations



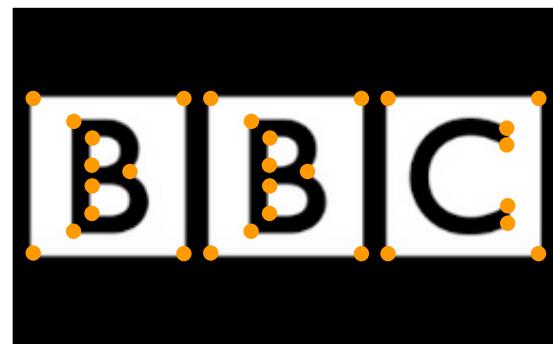
Harris Corner Detection Example



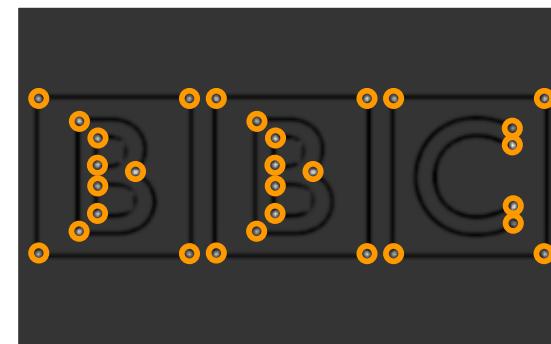
Image



Corner Response R

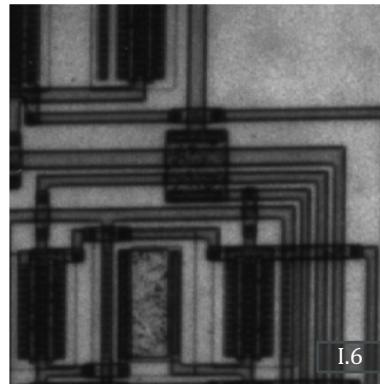


Detected Corners



Corner Response $R > T$

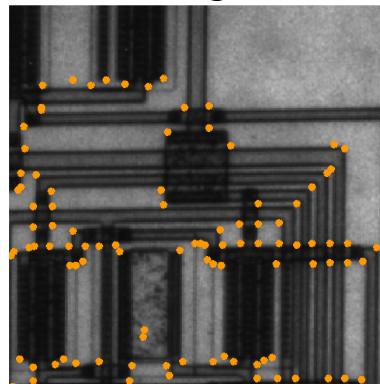
Harris Corner Detection Example (cont.)



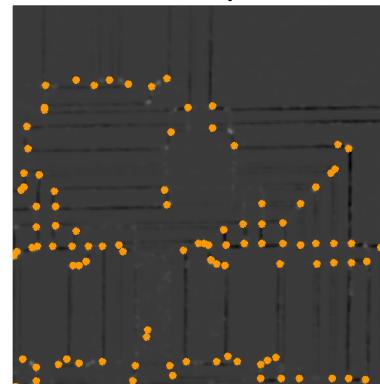
Image



Corner Response R



Detected Corners



Corner Response $R > T$

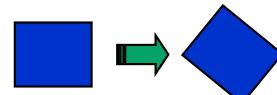
Invariance and Covariance

- We want features to be *invariant* to photometric transformations and *covariant* to geometric transformations
 - **Invariance:** image is transformed, and features do not change
 - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

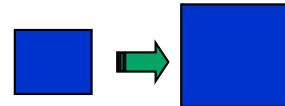


Transformations

- Geometric
 - **Rotation**



- Scale**

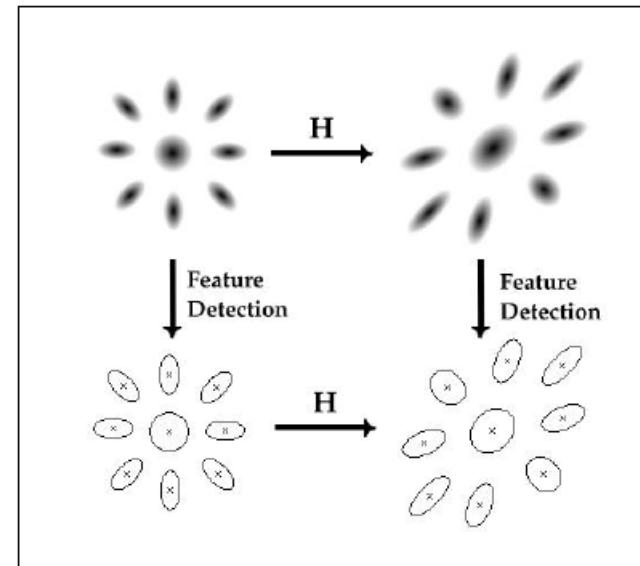
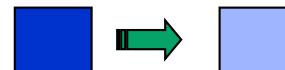


- **Affine**

valid for:
orthographic camera,
locally planar object

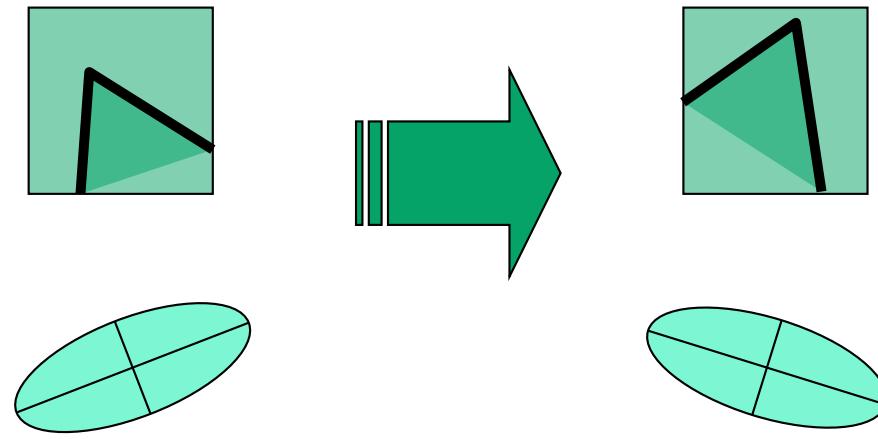


- Photometric
 - **Affine intensity change** ($I \rightarrow aI + b$)



T. Kadir, A. Zisserman, and M. Brady, An Affine Invariant Salient Region Detector, ECCV 2004

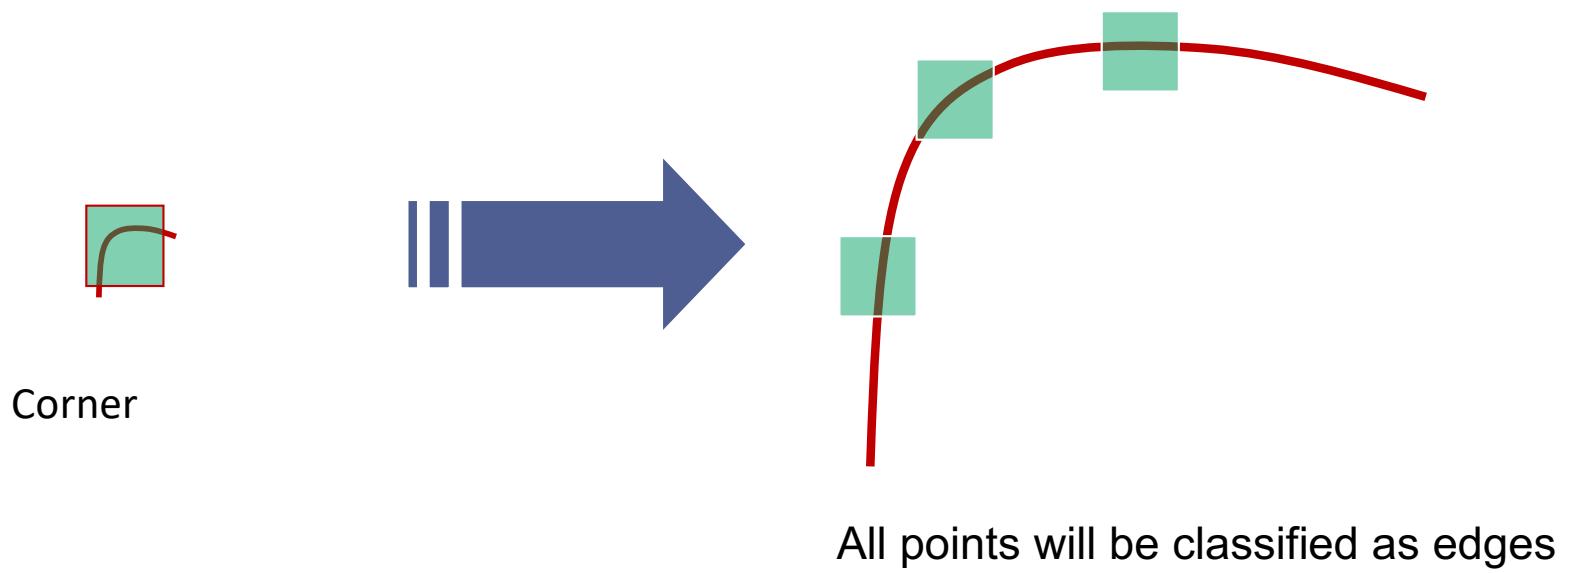
Image Rotation



Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response R is invariant w.r.t. rotation and corner location is covariant

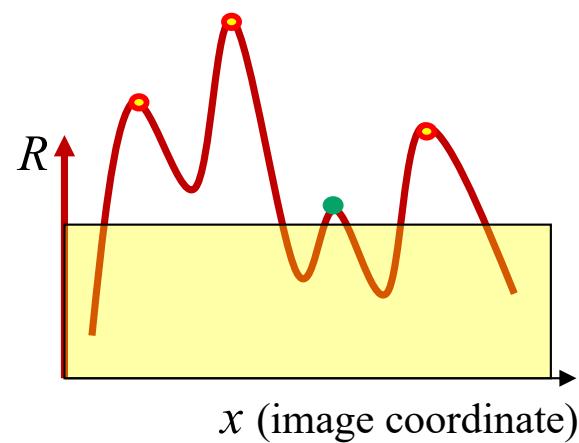
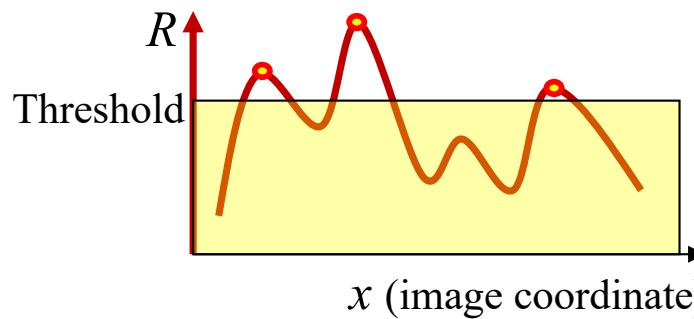
Scaling



Not invariant to scaling

Affine Intensity Change

- Only derivatives are used => invariance to intensity shift $I \rightarrow I + b$
- Intensity scale: $I \rightarrow a I$



Partially invariant to affine intensity change

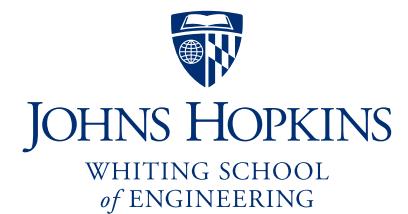
Summary

- **Corners:** Are one way to think about stable feature locations in images
- Essential concepts in this lecture:
 - How corners are expressed in real images
 - Picking corner points from the cornerness map
 - The idea of invariance and covariance

Johns Hopkins Engineering

Computer Vision

Corner Detection and Matching



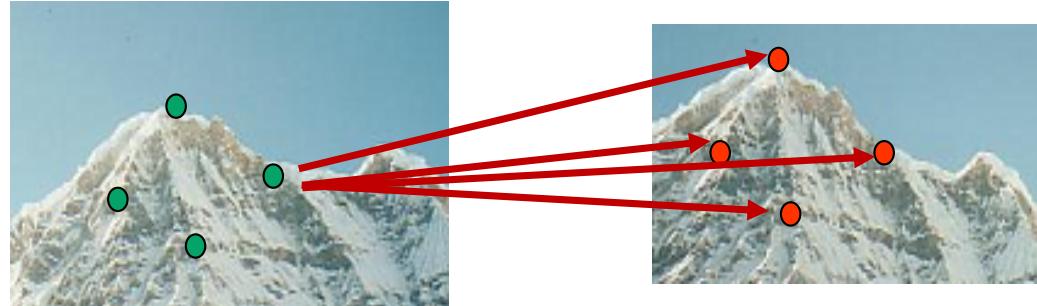
Corner Detection and Matching

- Topics:
 - Matching With Patches As Features

Matching With Features

- Problem 2:
 - For each point correctly recognize the corresponding one

?

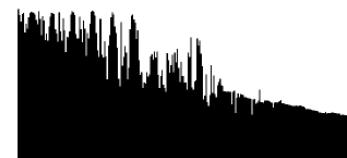
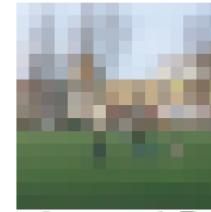


We need a reliable and distinctive descriptor

Cross-Correlation

- $$CC(P_1, P_2) = \frac{1}{N} \sum_i^N P_1[i]P_2[i]$$
- Output value depends on contrast and brightness of pixels
- Not invariant to changes in a, b

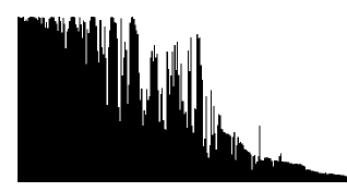
Affine photometric transformation:
 $I \rightarrow a I + b$



Original Patch and Intensity Values



Brightness Decreased, $CC = 0.262720397078039$



Contrast increased, $CC = 0.380413705374859$

Normalized Cross-Correlation

- Make each patch zero mean:

$$\mu = \frac{1}{N} \sum_{x,y} I(x,y)$$

$$Z(x,y) = I(x,y) - \mu$$

- Then make unit variance:

$$\sigma^2 = \frac{1}{N} \sum_{x,y} Z(x,y)^2$$

$$ZN(x,y) = \frac{Z(x,y)}{\sigma}$$

Affine photometric transformation:
 $I \rightarrow aI + b$



Original Patch and Intensity Values



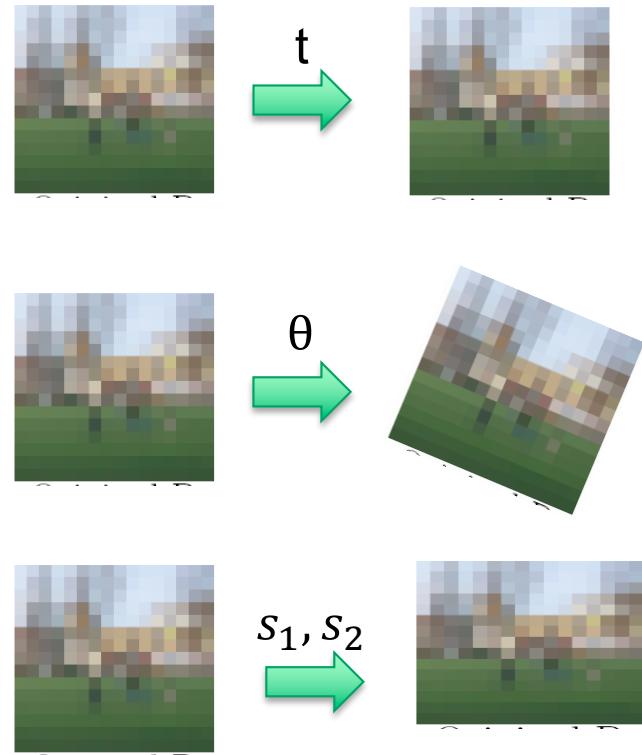
Brightness Decreased, CC = 0.99998895629



Contrast increased, CC = 0.969868160814

Normalized Cross-Correlation: Geometric Change

- Translation: Covariant
- Rotation
 - Not invariant
 - Can use eigenvectors to rotate patches
- Scale/Aspect ratio
 - Not invariant
 - Could be made invariant if there were a way to measure these changes in the detector



Feature Matching

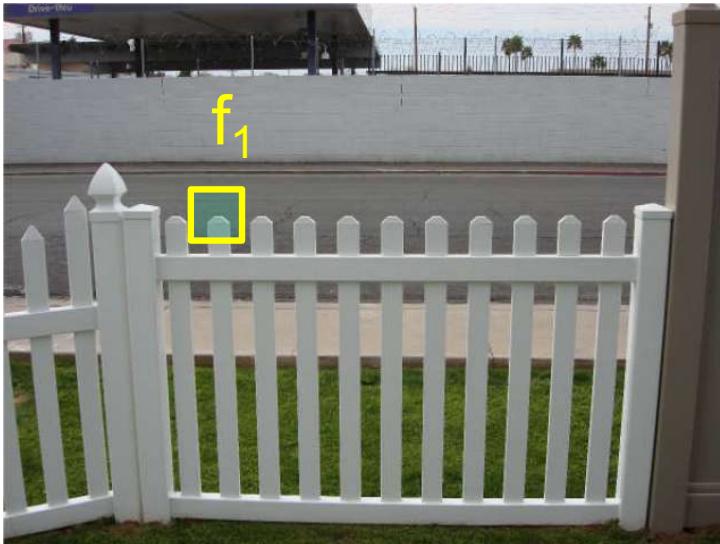
Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

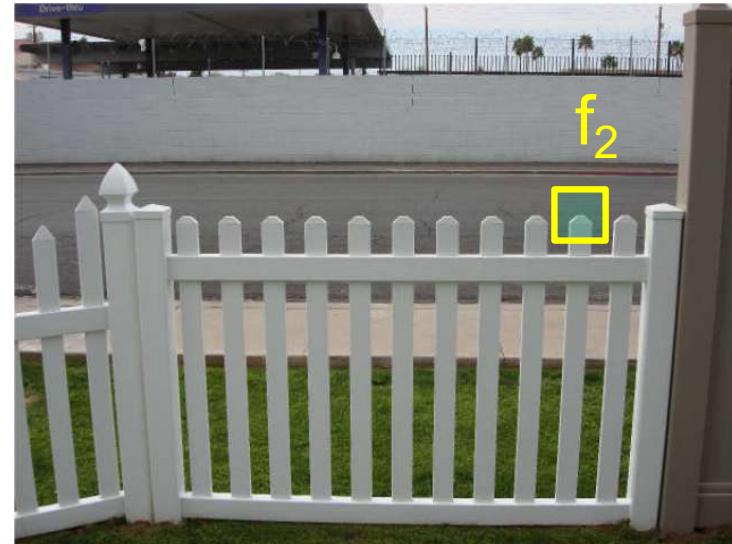
Feature Distance

How to define the difference between two features f_1, f_2 ?

- Simple approach is $\text{CC}(f_1, f_2)$
 - Cross-Correlation between entries of the two descriptors
 - Can give good scores to very ambiguous (bad) matches



I_1

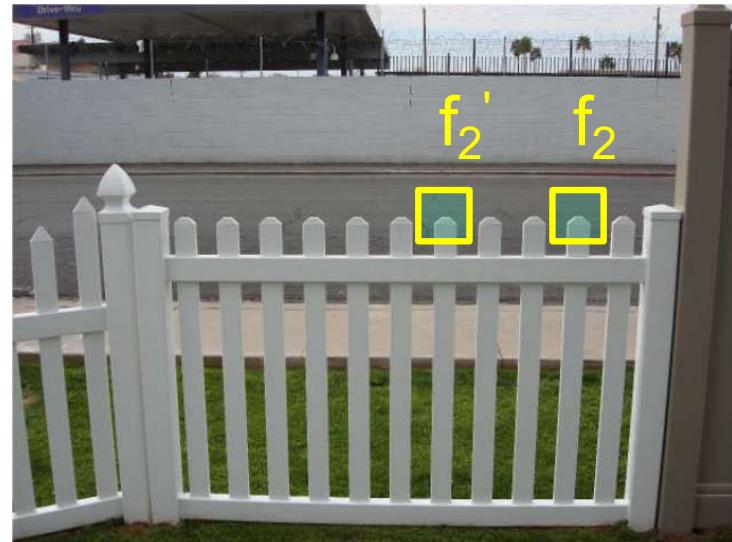
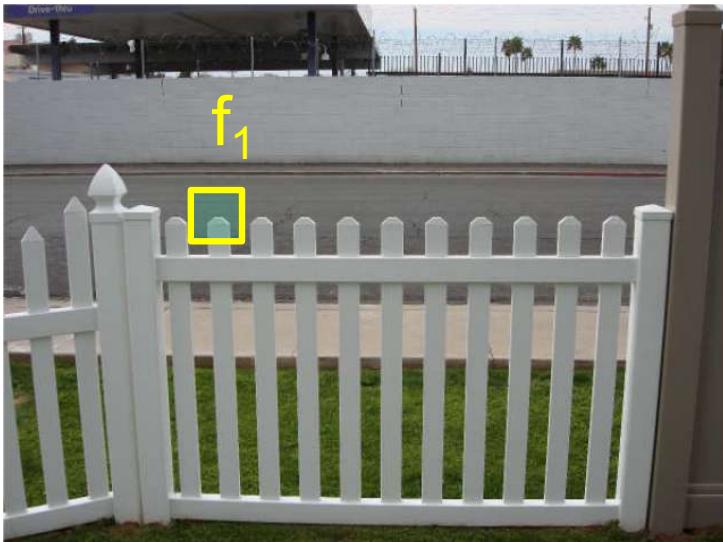


I_2

Feature Distance

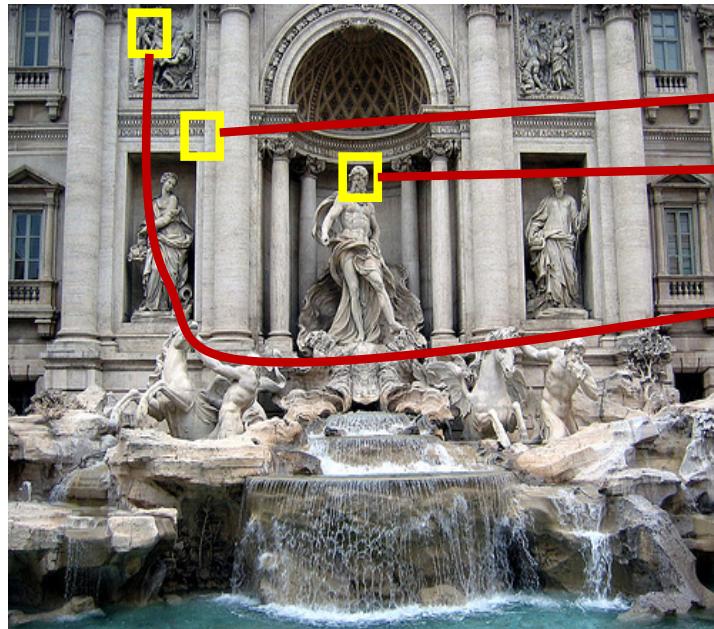
How to define the difference between two features f_1, f_2 ?

- Better approach: ratio distance = $CC(f_1, f_2) / CC(f_1, f'_2)$
 - f_2 is best CC match to f_1 in I_2
 - f'_2 is 2nd best CC match to f_1 in I_2
 - gives large values for ambiguous matches

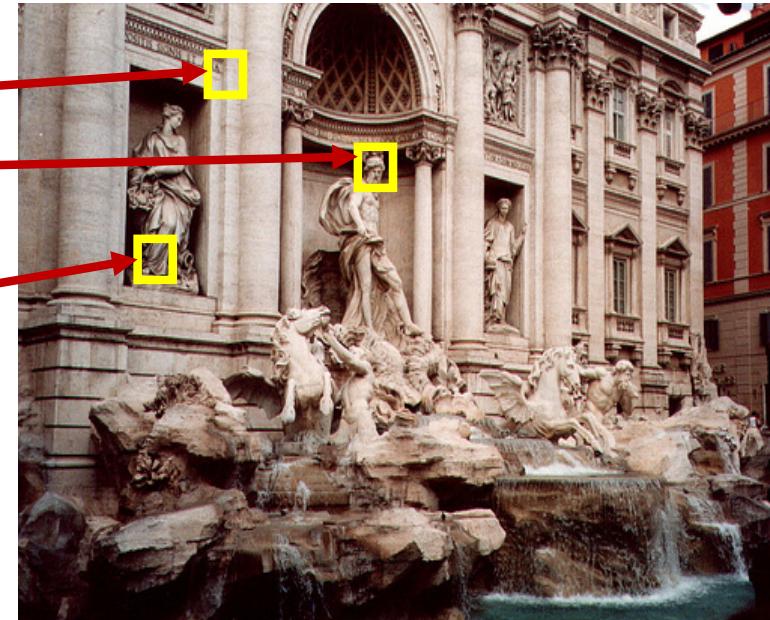


Evaluating The Results

How can we measure the performance of a feature matcher?

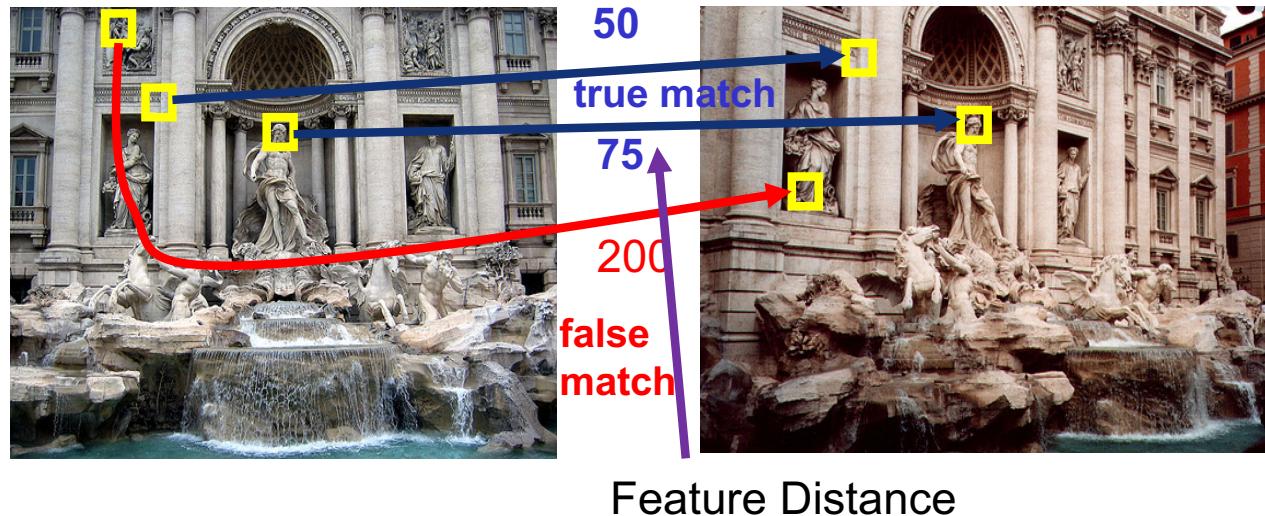


50
75
200



Feature Distance

True/False Positives



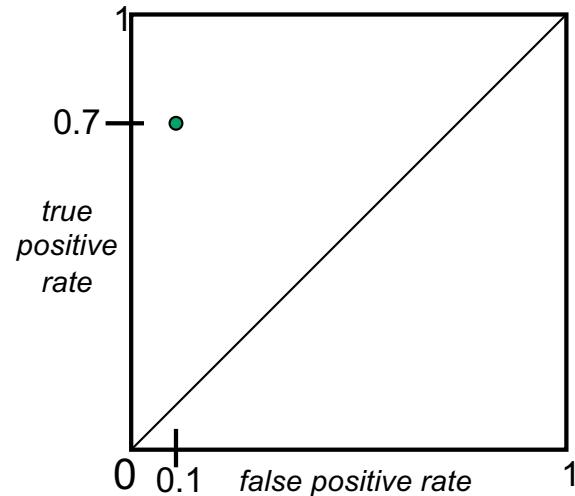
The distance threshold affects performance

- True positives = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Evaluating The Results

How can we measure the performance of a feature matcher?

$$\frac{\# \text{ true positives}}{\# \text{ matching features} \text{ (positives)}}$$



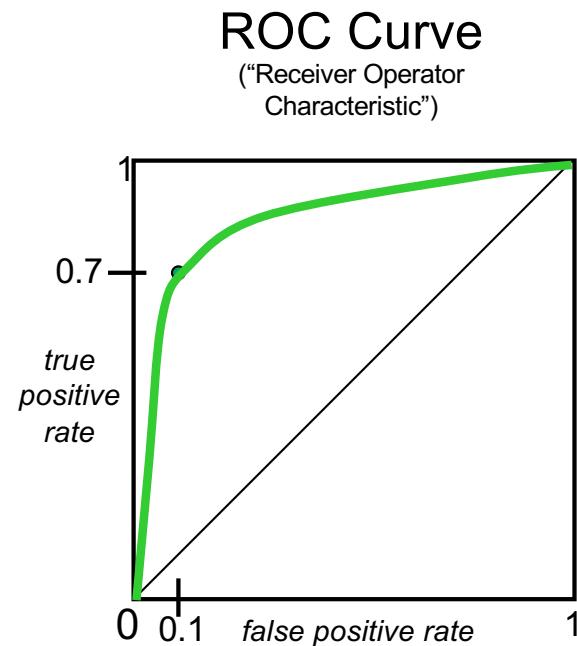
$$\frac{\# \text{ false positives}}{\# \text{ unmatched features} \text{ (negatives)}}$$

Evaluating The Results

How can we measure the performance of a feature matcher?

$$\frac{\# \text{ true positives}}{\# \text{ matching features} \text{ (positives)}}$$

- ROC Curves
 - Generated by counting # current/incorrect matches, for different thresholds
 - Want to maximize area under the curve (AUC)
 - Useful for comparing different feature matching methods



$$\frac{\# \text{ false positives}}{\# \text{ unmatched features} \text{ (negatives)}}$$

Summary

- **Corners:** Are one way to think about stable feature locations in images
- Essential concepts in this lecture:
 - Using correlation as matching – pros and cons
 - Tabulating matching results

References

- Textbooks:
 - Robot Vision (Chapter 8) Horn, B. K. P., MIT Press
 - Computer Vision: Algorithms and Applications (Chapter 4.1) Szelinski, 2011 (available online)
 - Computer Vision: A Modern Approach (Chapter 8) Forsyth, D and Ponce, J., Prentice Hall
 - Digital Image Processing (Chapter 3) González, R and Woods, R., Prentice Hall
- Papers:
 - [Canny1986] Canny, J., A Computational Approach To Edge Detection, IEEE Trans. Pattern Analysis and Machine Intelligence, 8(6):679–698, 1986.
 - [Harris1988] Harris, C. and Stephens, M., A combined corner and edge detector. Proceedings of the 4th Alvey Vision Conference. pp. 147–151.
 - [Marr1980] Marr, D. and Hildreth, E., Theory of Edge Detection,” Proc. R. Soc. London,B 207, 187-217, 1980.

Image Credits

- I.1 Adapted from Fig 3.1, Nalwa, V., A Guided Tour of Computer Vision.
- I.2 Adapted from Fig 3.3, Nalwa, V., A Guided Tour of Computer Vision.
- I.3 Matlab Demo Image
- I.4 http://en.wikipedia.org/wiki/File:Caf%C3%A9_wall.svg
- I.5 http://www.michaelbach.de/ot/geom_KitaokaBulge/index.html
- I.6 Matlab Demo Image



JOHNS HOPKINS
WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2020, All Rights Reserved.